

# 1) Data Science Skills [LinkedIn SQL Interview Question]

Given a table of candidates and their skills, you're tasked with finding the candidates best suited for an open Data Science job. You want to find candidates who are proficient in Python, Tableau, and PostgreSQL.

Write a query to list the candidates who possess all of the required skills for the job. Sort the output by candidate ID in ascending order.

## Assumption:

- There are no duplicates in the `candidates` table.

## `candidates` Table:

Column Name	Type
candidate_id	integer
skill	varchar

## `candidates` Example Input:

candidate_id	skill
123	Python
123	Tableau
123	PostgreSQL
234	R
234	PowerBI
234	SQL Server
345	Python
345	Tableau

## Example Output:

candidate_id
123

### **ANSWER**

```
SELECT
CANDIDATE_ID
FROM
CANDIDATES
WHERE
SKILL IN ('Python','Tableau','PostgreSQL')
GROUP BY
CANDIDATE_ID
HAVING
COUNT(CANDIDATE_ID) = 3
ORDER BY
CANDIDATE_ID ASC
```

## **2) Page With No Likes [Facebook SQL Interview Question]**

Assume you're given two tables containing data about Facebook Pages and their respective likes (as in "Like a Facebook Page").

Write a query to return the IDs of the Facebook pages that have zero likes. The output should be sorted in ascending order based on the page IDs.

**pages Table:**

Column Name	Type
page_id	integer
page_name	varchar

**Example Input:**

page_id	page_name
20001	SQL Solutions
20045	Brain Exercises
20701	Tips for Data Analysts

**page\_likes Table:**

Column Name	Type
user_id	integer
page_id	integer
liked_date	datetime

**page\_likes Example Input:**

user_id	page_id	liked_date
111	20001	04/08/2022 00:00:00
121	20045	03/12/2022 00:00:00
156	20001	07/25/2022 00:00:00

**Example Output:**

page_id
20701

**ANSWER**

```
SELECT
PAGE_ID
FROM
PAGES
EXCEPT
SELECT
PAGE_ID
```

FROM  
PAGE\_LIKES

### 3) Unfinished Parts [Tesla SQL Interview Question]

Tesla is investigating production bottlenecks and they need your help to extract the relevant data. Write a query to determine which parts have begun the assembly process but are not yet finished.

Assumptions:

- `parts_assembly` table contains all parts currently in production, each at varying stages of the assembly process.
- An unfinished part is one that lacks a `finish_date`.

This question is straightforward, so let's approach it with simplicity in both thinking and solution.

*Effective April 11th 2023, the problem statement and assumptions were updated to enhance clarity.*

#### `parts_assembly` Table

Column Name	Type
part	string
finish_date	datetime
assembly_step	integer

#### `parts_assembly` Example Input

part	finish_date	assembly_step
battery	01/22/2022 00:00:00	1
battery	02/22/2022 00:00:00	2
battery	03/22/2022 00:00:00	3
bumper	01/22/2022 00:00:00	1
bumper	02/22/2022 00:00:00	2

bumper		3
bumper		4

## Example Output

part	assembly_step
bumper	3
bumper	4

## Explanation

The bumpers in step 3 and 4 are the only item that remains unfinished as it lacks a recorded finish date.

The dataset you are querying against may have different input & output - **this is just an example!**

## ANSWER

```
SELECT
DISTINCT PART,
ASSEMBLY_STEP
FROM parts_assembly
WHERE FINISH_DATE IS NULL
```

## 4) Laptop vs. Mobile Viewership [New York Times SQL Interview Question]

This is the same question as problem #3 in the SQL Chapter of [Ace the Data Science Interview!](#)

Assume you're given the table on user viewership categorised by device type where the three types are laptop, tablet, and phone.

Write a query that calculates the total viewership for laptops and mobile devices where mobile is defined as the sum of tablet and phone viewership. Output the total viewership for laptops as `laptop_reviews` and the total viewership for mobile devices as `mobile_views`.

*Effective 15 April 2023, the solution has been updated with a more concise and easy-to-understand approach.*

## viewership Table

Column Name	Type
user_id	integer
device_type	string ('laptop', 'tablet', 'phone')
view_time	timestamp

## viewership Example Input

user_id	device_type	view_time
123	tablet	01/02/2022 00:00:00
125	laptop	01/07/2022 00:00:00
128	laptop	02/09/2022 00:00:00
129	phone	02/09/2022 00:00:00
145	tablet	02/24/2022 00:00:00

## Example Output

laptop_views	mobile_views
2	3

## Explanation

Based on the example input, there are a total of 2 laptop views and 3 mobile views.

The dataset you are querying against may have different input & output - **this is just an example!**

## ANSWER

```
SELECT
SUM(CASE WHEN DEVICE_TYPE = 'laptop' THEN 1 ELSE 0 END) AS LAPTOP_VIEWS,
SUM(CASE WHEN DEVICE_TYPE = 'phone' OR DEVICE_TYPE = 'tablet' THEN 1 ELSE 0
END) AS MOBILE_VIEWS
FROM
```

## 5) Histogram of Tweets [Twitter SQL Interview Question]

Assume you're given a table Twitter tweet data, write a query to obtain a histogram of tweets posted per user in 2022. Output the tweet count per user as the bucket and the number of Twitter users who fall into that bucket.

In other words, group the users by the number of tweets they posted in 2022 and count the number of users in each group.

`tweets` **Table:**

Column Name	Type
tweet_id	integer
user_id	integer
msg	string
tweet_date	timestamp

`tweets` **Example Input:**

tweet_id	user_id	msg	tweet_date
214252	111	Am considering taking Tesla private at \$420. Funding secured.	12/30/2021 00:00:00
739252	111	Despite the constant negative press covfefe	01/01/2022 00:00:00
846402	111	Following @NickSinghTech on Twitter changed my life!	02/14/2022 00:00:00
241425	254	If the salary is so competitive why won't you tell me what it is?	03/01/2022 00:00:00

231574	148	I no longer have a manager. I can't be managed	03/23/2022 00:00:00
--------	-----	--	------------------------

### Example Output:

tweet_bucket	users_num
1	2
2	1

### Explanation:

Based on the example output, there are two users who posted only one tweet in 2022, and one user who posted two tweets in 2022. The query groups the users by the number of tweets they posted and displays the number of users in each group.

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```

SELECT
TWEET_COUNT_BY_USER AS TWEET_BUCKET,
COUNT(USER_ID) AS USERS_NUM
FROM
(SELECT
  USER_ID,
  COUNT(TWEET_ID) AS TWEET_COUNT_BY_USER
FROM
  TWEETS
WHERE
  TWEET_DATE BETWEEN '01/01/2022' AND '12/31/2022'
GROUP BY
  USER_ID) AS TEMP_TABLE
GROUP BY TWEET_BUCKET

```

### **ALTERNATE ANSWER (USING CTE)**

```

WITH total_tweets AS (
  SELECT
    user_id,
    COUNT(tweet_id) AS tweet_count_per_user
  FROM tweets

```



```
WHERE tweet_date BETWEEN '2022-01-01'
AND '2022-12-31'
GROUP BY user_id)
```

```
SELECT
tweet_count_per_user AS tweet_bucket,
COUNT(user_id) AS users_num
FROM total_tweets
GROUP BY tweet_count_per_user;
```

## 6) Average Post Hiatus (Part 1)

### [Facebook SQL Interview Question]

Given a table of Facebook posts, for each user who posted at least twice in 2021, write a query to find the number of days between each user's first post of the year and last post of the year in the year 2021. Output the user and number of the days between each user's first and last post.

p.s. If you've read the [Ace the Data Science Interview](#) and liked it, consider writing us a review?

**posts Table:**

Column Name	Type
user_id	integer
post_id	integer
post_date	timestamp
post_content	text

**posts Example Input:**

user_id	post_id	post_date	post_content
151652	599415	07/10/2021 12:00:00	Need a hug

661093	624356	07/29/2021 13:00:00	Bed. Class 8-12. Work 12-3. Gym 3-5 or 6. Then class 6-10. Another day that's gonna fly by. I miss my girlfriend
004239	784254	07/04/2021 11:00:00	Happy 4th of July!
661093	442560	07/08/2021 14:00:00	Just going to cry myself to sleep after watching Marley and Me.
151652	111766	07/12/2021 19:00:00	I'm so done with covid - need travelling ASAP!

### Example Output:

user_id	days_between
151652	2
661093	21

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```

SELECT
USER_ID,
MAX(POST_DATE::DATE) - MIN(POST_DATE::DATE) AS DAYS_BETWEEN
FROM
POSTS
WHERE
POST_DATE BETWEEN '01/01/2021' AND '12/31/2021'
GROUP BY
USER_ID
HAVING
COUNT(USER_ID) >= 2

```

## 7) Teams Power Users [Microsoft SQL Interview Question]

Write a query to identify the top 2 Power Users who sent the highest number of messages on Microsoft Teams in August 2022. Display the IDs of these 2 users along with the total number of messages they sent. Output the results in descending order based on the count of the messages.

Assumption:

- No two users have sent the same number of messages in August 2022.

**messages Table:**

Column Name	Type
message_id	integer
sender_id	integer
receiver_id	integer
content	varchar
sent_date	datetime

**messages Example Input:**

message_id	sender_id	receiver_id	content	sent_date
901	3601	4500	You up?	08/03/2022 00:00:00
902	4500	3601	Only if you're buying	08/03/2022 00:00:00
743	3601	8752	Let's take this offline	06/14/2022 00:00:00
922	3601	4500	Get on the call	08/10/2022 00:00:00

**Example Output:**

sender_id	message_count
3601	2

4500	1
------	---

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```
SELECT
SENDER_ID,
COUNT(SENDER_ID) AS MESSAGE_COUNT
FROM
MESSAGES
WHERE
SENT_DATE BETWEEN '08/01/2022' AND '08/31/2022'
GROUP BY
SENDER_ID
ORDER BY
MESSAGE_COUNT DESC
LIMIT
2
```

## **8) Duplicate Job Listings [Linkedin SQL Interview Question]**

Assume you're given a table containing job postings from various companies on the LinkedIn platform. Write a query to retrieve the count of companies that have posted duplicate job listings.

Definition:

- Duplicate job listings are defined as two job listings within the same company that share identical titles and descriptions.

`job_listings` **Table:**

Column Name	Type
job_id	integer
company_id	integer
title	string

description	string
-------------	--------

### job\_listings Example Input:

job_id	company_id	title	description
248	827	Business Analyst	Business analyst evaluates past and current business data with the primary goal of improving decision-making processes within organizations.
149	845	Business Analyst	Business analyst evaluates past and current business data with the primary goal of improving decision-making processes within organizations.
945	345	Data Analyst	Data analyst reviews data to identify key insights into a business's customers and ways the data can be used to solve problems.
164	345	Data Analyst	Data analyst reviews data to identify key insights into a business's customers and ways the data can be used to solve problems.
172	244	Data Engineer	Data engineer works in a variety of settings to build systems that collect, manage, and convert raw data into usable information for data scientists and business analysts to interpret.

### Example Output:

duplicate_companies
1

### Explanation:

There is one company ID 345 that posted duplicate job listings. The duplicate listings, IDs 945 and 164 have identical titles and descriptions.

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```
SELECT
COUNT(DISTINCT COMPANY_ID) AS DUPLICATE_COMPANIES
FROM(
  SELECT
    COMPANY_ID,
    TITLE,
    DESCRIPTION,
    COUNT(JOB_ID) AS COUNTER
  FROM
    JOB_LISTINGS
  GROUP BY
    COMPANY_ID,
    TITLE,
    DESCRIPTION) AS TEMPORARY_TABLE
WHERE
COUNTER >= 2
```

## **9) Cities With Completed Trades [Robinhood SQL Interview Question]**

Assume you're given the tables containing completed trade orders and user details in a Robinhood trading system.

Write a query to retrieve the top three cities that have the highest number of completed trade orders listed in descending order. Output the city name and the corresponding number of completed trade orders.

**trades Table:**

Column Name	Type
order_id	integer
user_id	integer

price	decimal
quantity	integer
status	string('Completed' ,'Cancelled')
timestamp	datetime

### trades **Example Input:**

order_id	user_id	price	quantity	status	timestamp
100101	111	9.80	10	Cancelled	08/17/2022 12:00:00
100102	111	10.00	10	Completed	08/17/2022 12:00:00
100259	148	5.10	35	Completed	08/25/2022 12:00:00
100264	148	4.80	40	Completed	08/26/2022 12:00:00
100305	300	10.00	15	Completed	09/05/2022 12:00:00
100400	178	9.90	15	Completed	09/09/2022 12:00:00
100565	265	25.60	5	Completed	12/19/2022 12:00:00

### users **Table:**

Column Name	Type
user_id	integer
city	string
email	string
signup_date	datetime

**users** Example Input:

user_id	city	email	signup_date
111	San Francisco	rrrok10@gmail.com	08/03/2021 12:00:00
148	Boston	sailor9820@gmail.com	08/20/2021 12:00:00
178	San Francisco	harrypotterfan182@gmail.com	01/05/2022 12:00:00
265	Denver	shadower_@hotmail.com	02/26/2022 12:00:00
300	San Francisco	houstoncowboy1122@hotmail.com	06/30/2022 12:00:00

**Example Output:**

city	total_orders
San Francisco	3
Boston	2
Denver	1

## Explanation

In the given dataset, San Francisco has the highest number of completed trade orders with 3 orders. Boston holds the second position with 2 orders, and Denver ranks third with 1 order.

The dataset you are querying against may have different input & output - **this is just an example!**

## ANSWER

```
SELECT
USERS.CITY,
COUNT(TRADES.ORDER_ID) AS TOTAL_ORDERS
FROM
```



```
USERS INNER JOIN TRADES
  ON USERS.USER_ID = TRADES.USER_ID
WHERE
STATUS = 'Completed'
GROUP BY
USERS.CITY
ORDER BY
TOTAL_ORDERS DESC
LIMIT 3
```

## 10) Average Review Ratings [Amazon SQL Interview Question]

Given the reviews table, write a query to retrieve the average star rating for each product, grouped by month. The output should display the month as a numerical value, product ID, and average star rating rounded to two decimal places. Sort the output first by month and then by product ID.

P.S. If you've read the [Ace the Data Science Interview](#), and liked it, consider writing us a review?

**reviews Table:**

Column Name	Type
review_id	integer
user_id	integer
submit_date	datetime
product_id	integer
stars	integer (1-5)

**reviews Example Input:**

review_id	user_id	submit_date	product_id	stars
6171	123	06/08/2022 00:00:00	50001	4
7802	265	06/10/2022 00:00:00	69852	4

5293	362	06/18/2022 00:00:00	50001	3
6352	192	07/26/2022 00:00:00	69852	3
4517	981	07/05/2022 00:00:00	69852	2

### Example Output:

mth	product	avg_stars
6	50001	3.50
6	69852	4.00
7	69852	2.50

### Explanation

Product 50001 received two ratings of 4 and 3 in the month of June (6th month), resulting in an average star rating of 3.5.

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```
SELECT
EXTRACT(MONTH FROM submit_date) AS MTH,
PRODUCT_ID AS PRODUCT,
ROUND(AVG(STARS),2) AS AVG_STARS
FROM
REVIEWS
GROUP BY
MTH,PRODUCT
ORDER BY
MTH,PRODUCT
```

## 11) App Click-through Rate (CTR) [Facebook SQL Interview Question]

Assume you have an events table on Facebook app analytics. Write a query to calculate the click-through rate (CTR) for the app in 2022 and round the results to 2 decimal places.

Definition and note:

- Percentage of click-through rate (CTR) =  $100.0 * \text{Number of clicks} / \text{Number of impressions}$
- To avoid integer division, multiply the CTR by 100.0, not 100.

`events` **Table:**

Column Name	Type
app_id	integer
event_type	string
timestamp	datetime

`events` **Example Input:**

app_id	event_type	timestamp
123	impression	07/18/2022 11:36:12
123	impression	07/18/2022 11:37:12
123	click	07/18/2022 11:37:42
234	impression	07/18/2022 14:15:12
234	click	07/18/2022 14:16:12

**Example Output:**

app_id	ctr
123	50.00
234	100.00

## Explanation

Let's consider an example of App 123. This app has a click-through rate (CTR) of 50.00% because out of the 2 impressions it received, it got 1 click.

To calculate the CTR, we divide the number of clicks by the number of impressions, and then multiply the result by 100.0 to express it as a percentage. In this case, 1 divided by 2 equals 0.5, and when multiplied by 100.0, it becomes 50.00%. So, the CTR of App 123 is 50.00%.

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```
SELECT
APP_ID,
ROUND(
100.0*
SUM(CASE WHEN EVENT_TYPE = 'click' THEN 1 ELSE 0 END)/
SUM(CASE WHEN EVENT_TYPE = 'impression' THEN 1 ELSE 0 END), 2) AS CTR
FROM
EVENTS
WHERE
TIMESTAMP BETWEEN '01/01/2022' AND '12/31/2022'
GROUP BY
APP_ID
```

## **12) Second Day Confirmation [TikTok SQL Interview Question]**

Assume you're given tables with information about TikTok user sign-ups and confirmations through email and text. New users on TikTok sign up using their email addresses, and upon sign-up, each user receives a text message confirmation to activate their account.

Write a query to display the user IDs of those who did not confirm their sign-up on the first day, but confirmed on the second day.

Definition:

- `action_date` refers to the date when users activated their accounts and confirmed their sign-up through text messages.

`emails` **Table:**

Column Name	Type
email_id	integer
user_id	integer
signup_date	datetime

**emails** Example Input:

email_id	user_id	signup_date
125	7771	06/14/2022 00:00:00
433	1052	07/09/2022 00:00:00

**texts** Table:

Column Name	Type
text_id	integer
email_id	integer
signup_action	string ('Confirmed', 'Not confirmed')
action_date	datetime

**texts** Example Input:

text_id	email_id	signup_action	action_date
6878	125	Confirmed	06/14/2022 00:00:00
6997	433	Not Confirmed	07/09/2022 00:00:00
7000	433	Confirmed	07/10/2022 00:00:00

**Example Output:**

user_id
1052

## Explanation:

Only User 1052 confirmed their sign-up on the second day.

The dataset you are querying against may have different input & output - **this is just an example!**

## ANSWER

```
SELECT
EMAILS.USER_ID
FROM
EMAILS INNER JOIN TEXTS
ON EMAILS.EMAIL_ID = TEXTS.EMAIL_ID
WHERE
TEXTS.ACTION_DATE = EMAILS.SIGNUP_DATE + INTERVAL '1 day'
AND TEXTS.SIGNUP_ACTION = 'Confirmed'
```

# 13) Cards Issued Difference [JPMorgan Chase SQL Interview Question]

Your team at JPMorgan Chase is preparing to launch a new credit card, and to gain some insights, you're analyzing how many credit cards were issued each month.

Write a query that outputs the name of each credit card and the difference in the number of issued cards between the month with the highest issuance cards and the lowest issuance. Arrange the results based on the largest disparity.

## monthly\_cards\_issued Table:

Column Name	Type
issue_month	integer
issue_year	integer
card_name	string
issued_amount	integer

## monthly\_cards\_issued Example Input:

card_name	issued_amount	issue_month	issue_year
-----------	---------------	-------------	------------

Chase Freedom Flex	55000	1	2021
Chase Freedom Flex	60000	2	2021
Chase Freedom Flex	65000	3	2021
Chase Freedom Flex	70000	4	2021
Chase Sapphire Reserve	170000	1	2021
Chase Sapphire Reserve	175000	2	2021
Chase Sapphire Reserve	180000	3	2021

### Example Output:

card_name	difference
Chase Freedom Flex	15000
Chase Sapphire Reserve	10000

### Explanation:

Chase Freedom Flex's best month was 70k cards issued and the worst month was 55k cards, so the difference is 15k cards.

Chase Sapphire Reserve's best month was 180k cards issued and the worst month was 170k cards, so the difference is 10k cards.

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```
SELECT
DISTINCT CARD_NAME,
MAX(ISSUED_AMOUNT) - MIN(ISSUED_AMOUNT) AS DIFFERENCE
FROM
MONTHLY_CARDS_ISSUED
GROUP BY
CARD_NAME
```

## 14) Compressed Mean [Alibaba SQL Interview Question]

You're trying to find the mean number of items per order on Alibaba, rounded to 1 decimal place using tables which includes information on the count of items in each order (`item_count` table) and the corresponding number of orders for each item count (`order_occurrences` table).

`items_per_order` Table:

Column Name	Type
item_count	integer
order_occurrences	integer

`items_per_order` Example Input:

item_count	order_occurrences
1	500
2	1000
3	800
4	1000

There are a total of 500 orders with one item per order, 1000 orders with two items per order, and 800 orders with three items per order."

Example Output:

mean
2.7

### Explanation

Let's calculate the arithmetic average:

Total items =  $(1 \times 500) + (2 \times 1000) + (3 \times 800) + (4 \times 1000) = 8900$



Total orders = 500 + 1000 + 800 + 1000 = 3300

Mean = 8900 / 3300 = 2.7

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```
SELECT  
ROUND(SUM(ITEM_COUNT :: DECIMAL * ORDER_OCCURRENCES) /  
SUM(ORDER_OCCURRENCES),1) AS SID  
FROM  
ITEMS_PER_ORDER
```

## **15) Pharmacy Analytics (Part 1) [CVS Health SQL Interview Question]**

CVS Health is trying to better understand its pharmacy sales, and how well different products are selling. Each drug can only be produced by one manufacturer.

Write a query to find the top 3 most profitable drugs sold, and how much profit they made. Assume that there are no ties in the profits. Display the result from the highest to the lowest total profit.

### **Definition:**

- `cogs` stands for Cost of Goods Sold which is the direct cost associated with producing the drug.
- Total Profit = Total Sales - Cost of Goods Sold

If you like this question, try out [Pharmacy Analytics \(Part 2\)](#)!

### **pharmacy\_sales Table:**

Column Name	Type
product_id	integer
units_sold	integer
total_sales	decimal

cogs	decimal
manufacturer	varchar
drug	varchar

pharmacy\_sales **Example Input:**

product_id	units_sold	total_sales	cogs	manufacturer	drug
9	37410	293452.54	208876.01	Eli Lilly	Zyprexa
34	94698	600997.19	521182.16	AstraZeneca	Surmontil
61	77023	500101.61	419174.97	Biogen	Varicose Relief
136	144814	1084258	1006447.73	Biogen	Burkhart

**Example Output:**

drug	total_profit
Zyprexa	84576.53
Varicose Relief	80926.64
Surmontil	79815.03

**Explanation:**

Zyprexa made the most profit (of \$84,576.53) followed by Varicose Relief (of \$80,926.64) and Surmontil (of \$79,815.3).

The dataset you are querying against may have different input & output - **this is just an example!**

**ANSWER**

```
SELECT
DRUG,
TOTAL_SALES - COGS AS TOTAL_PROFIT
FROM
PHARMACY_SALES
```

```
ORDER BY
TOTAL_PROFIT DESC
LIMIT 3
```

## 16) Pharmacy Analytics (Part 2) [CVS Health SQL Interview Question]

CVS Health is analyzing its pharmacy sales data, and how well different products are selling in the market. Each drug is exclusively manufactured by a single manufacturer.

Write a query to identify the manufacturers associated with the drugs that resulted in losses for CVS Health and calculate the total amount of losses incurred.

Output the manufacturer's name, the number of drugs associated with losses, and the total losses in absolute value. Display the results sorted in descending order with the highest losses displayed at the top.

If you like this question, try out [Pharmacy Analytics \(Part 3\)](#)!

**pharmacy\_sales Table:**

Column Name	Type
product_id	integer
units_sold	integer
total_sales	decimal
cogs	decimal
manufacturer	varchar
drug	varchar

**pharmacy\_sales Example Input:**

product_id	units_sold	total_sales	cogs	manufacturer	drug
156	89514	3130097.00	3427421.73	Biogen	Acyclovir

25	222331	2753546.00	2974975.36	AbbVie	Lamivudine and Zidovudine
50	90484	2521023.73	2742445.90	Eli Lilly	Dermasorb TA Complete Kit
98	110746	813188.82	140422.87	Biogen	Medi-Chord

### Example Output:

manufacturer	drug_count	total_loss
Biogen	1	297324.73
AbbVie	1	221429.36
Eli Lilly	1	221422.17

### Explanation:

The first three rows indicate that some drugs resulted in losses. Among these, Biogen had the highest losses, followed by AbbVie and Eli Lilly. However, the Medi-Chord drug manufactured by Biogen reported a profit and was excluded from the result.

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```

SELECT
MANUFACTURER,
COUNT(DRUG) AS DRUG_COUNT,
ABS(SUM(TOTAL_SALES - COGS)) AS TOTAL_LOSS
FROM
PHARMACY_SALES
WHERE
TOTAL_SALES - COGS < 0
GROUP BY
MANUFACTURER

```

ORDER BY  
TOTAL\_LOSS DESC

## 17) Pharmacy Analytics (Part 3) [CVS Health SQL Interview Question]

CVS Health wants to gain a clearer understanding of its pharmacy sales and the performance of various products.

Write a query to calculate the total drug sales for each manufacturer. Round the answer to the nearest million and report your results in descending order of total sales. In case of any duplicates, sort them alphabetically by the manufacturer name.

Since this data will be displayed on a dashboard viewed by business stakeholders, please format your results as follows: "\$36 million".

If you like this question, try out [Pharmacy Analytics \(Part 4\)](#)!

**pharmacy\_sales Table:**

Column Name	Type
product_id	integer
units_sold	integer
total_sales	decimal
cogs	decimal
manufacturer	varchar
drug	varchar

**pharmacy\_sales Example Input:**

product_id	units_sold	total_sales	cogs	manufacturer	drug
94	132362	2041758.41	1373721.70	Biogen	UP and UP
9	37410	293452.54	208876.01	Eli Lilly	Zyprexa

50	90484	2521023.73	2742445.9	Eli Lilly	Dermasorb
61	77023	500101.61	419174.97	Biogen	Varicose Relief
136	144814	1084258.00	1006447.73	Biogen	Burkhart

### Example Output:

manufacturer	sale
Biogen	\$4 million
Eli Lilly	\$3 million

### Explanation

The total sales for Biogen is \$4 million ( $\$2,041,758.41 + \$500,101.61 + \$1,084,258.00 = \$3,626,118.02$ ) and for Eli Lilly is \$3 million ( $\$293,452.54 + \$2,521,023.73 = \$2,814,476.27$ ).

The dataset you are querying against may have different input & output - **this is just an example!**

### **ANSWER**

```
SELECT
MANUFACTURER,
CONCAT('$',ROUND(SUM(TOTAL_SALES)/1000000,0), ' million') AS SALES_MIL
FROM
PHARMACY_SALES
GROUP BY
MANUFACTURER
ORDER BY
ROUND(SUM(TOTAL_SALES)/1000000,0) DESC,
MANUFACTURER DESC
```

